



Information Technology Institute



Operating System Fundamentals

Chapter Four

PROCESSES

Table of Content

- Process Concept
- Process Scheduling
- Operations on Processes
- Threads
- Cooperating Processes
- Inter-process Communication

PROCESS CONCEPT

Process Concept

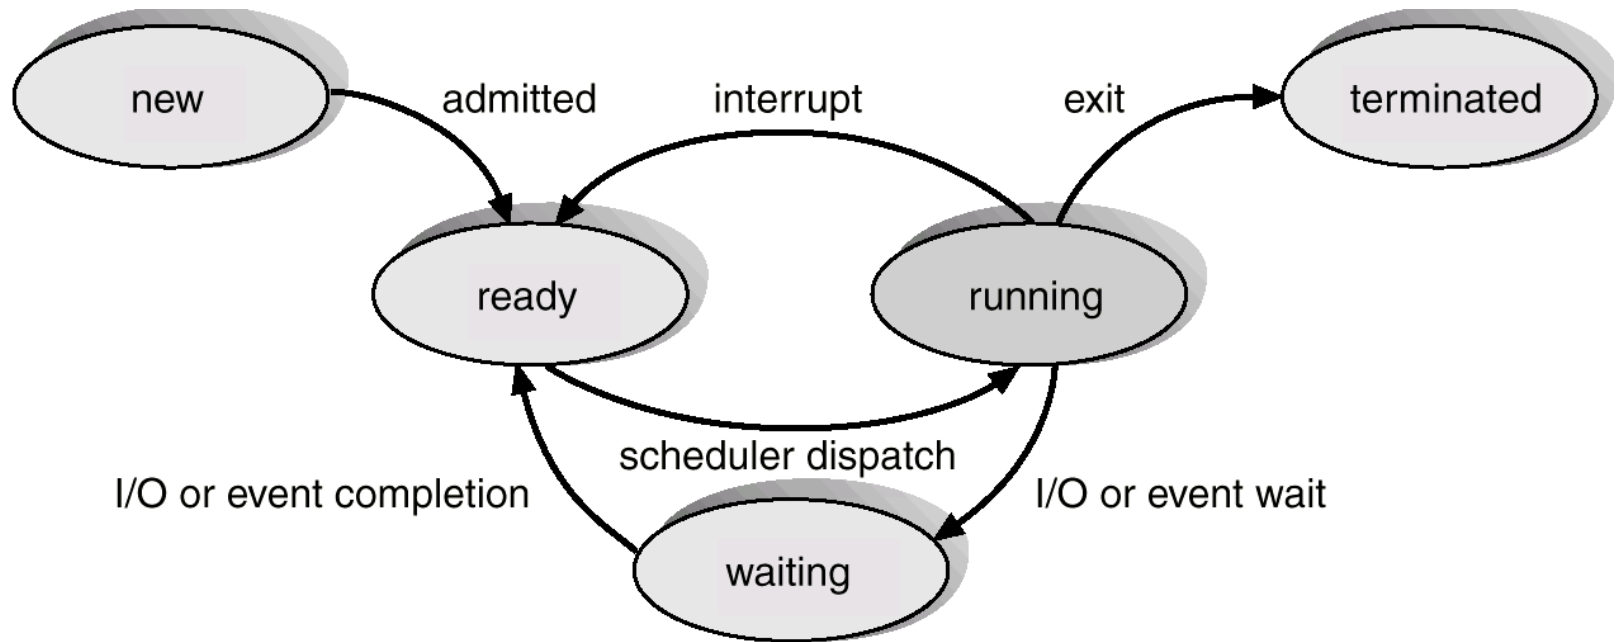
- Process – a program in execution; process execution must progress in sequential fashion.
- An operating system executes a variety of programs:
 - Batch system
 - Time-shared systems

* Textbook uses the terms job and process almost interchangeably

Process Contents

- **Text section**
 - Program instructions
- **Program counter**
 - Next instruction
- **Stack**
 - Local variables
 - Return addresses
 - Method parameters
- **Data section**
 - Global variables

Process State



Process State Cont'd

- **As a process executes, it changes state**
 - new: The process is being created.
 - running: Instructions are being executed.
 - waiting: The process is waiting for some event to occur.
 - ready: The process is waiting to be assigned to a processor.
 - terminated: The process has finished execution.

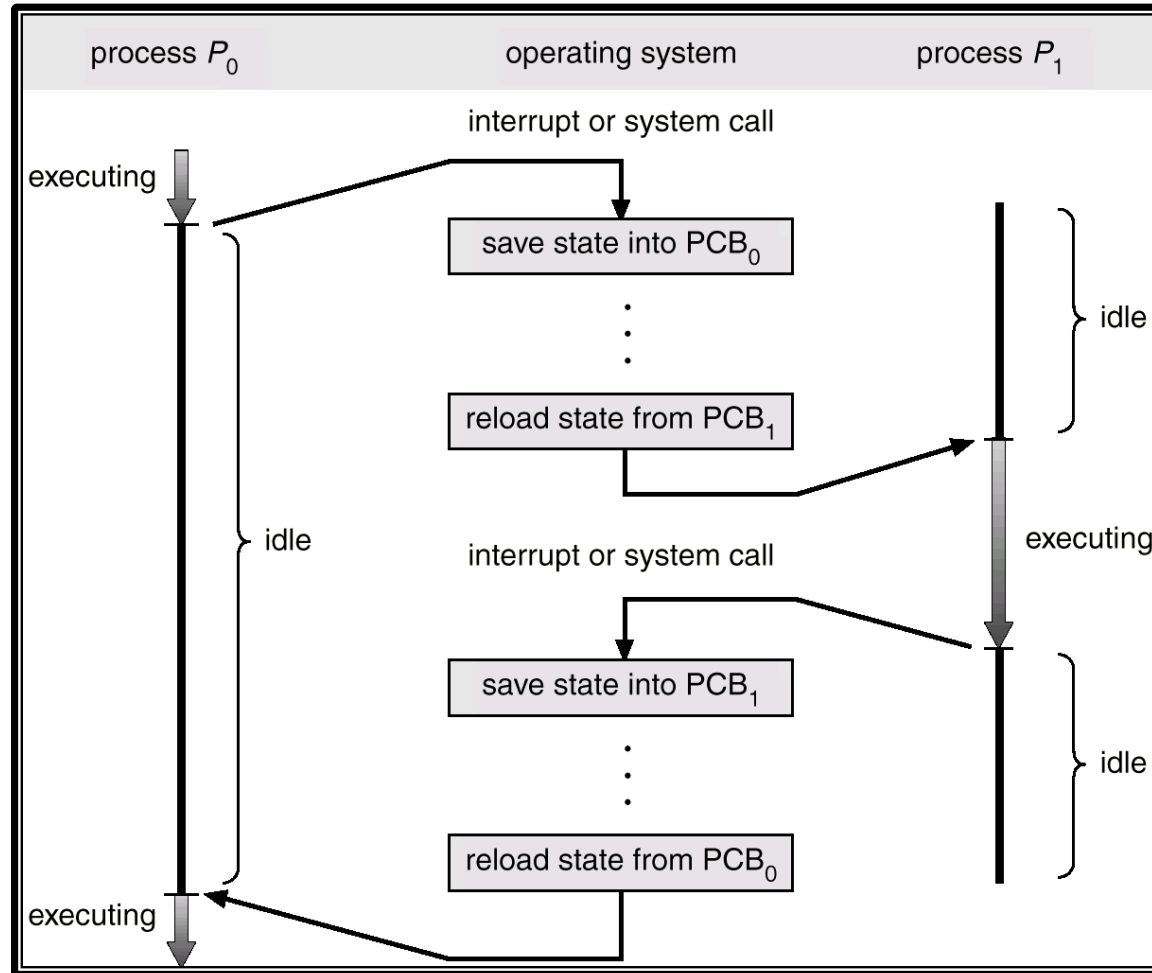
Process Control Block (PCB)

- Information associated with each process.
 - Process state
 - Program counter
 - CPU registers
 - CPU scheduling information
 - Memory-management information
 - Accounting information
 - I/O status information

Process Control Block (PCB) Cont'd

Pointer	Process state
Process number	
Program counter	
CPU registers	
Memory management info	
I/O status information	
Accounting Information	

CPU Switch From Process to Process



PROCESS SCHEDULING

Process Scheduling

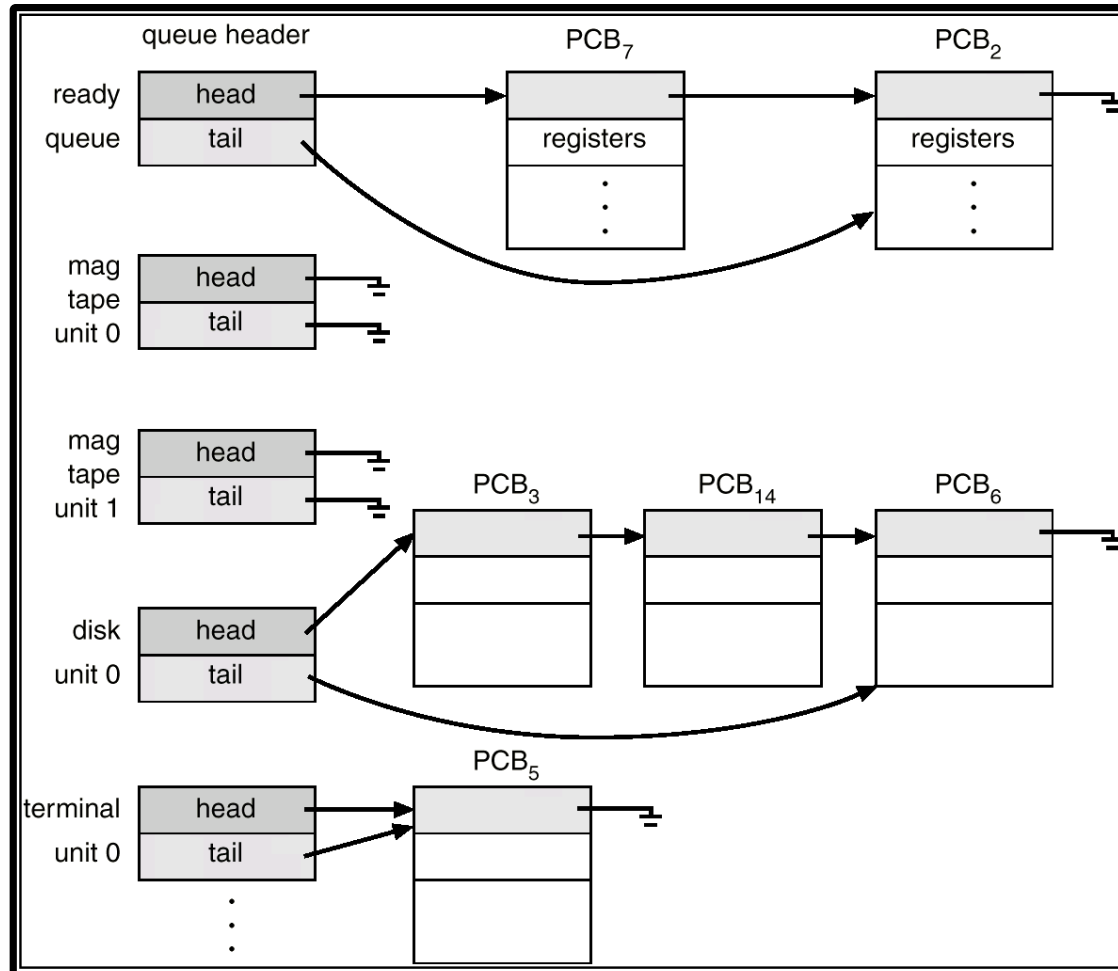
- **Multi-Programming systems**
 - Some processes executing at all times
 - Maximize CPU utilization
- **Time-Sharing systems**
 - Switch the CPU among processes frequently
 - Users can interact with a program while it's executing
 - Virtually run processes at the same time

Process Scheduling Queues

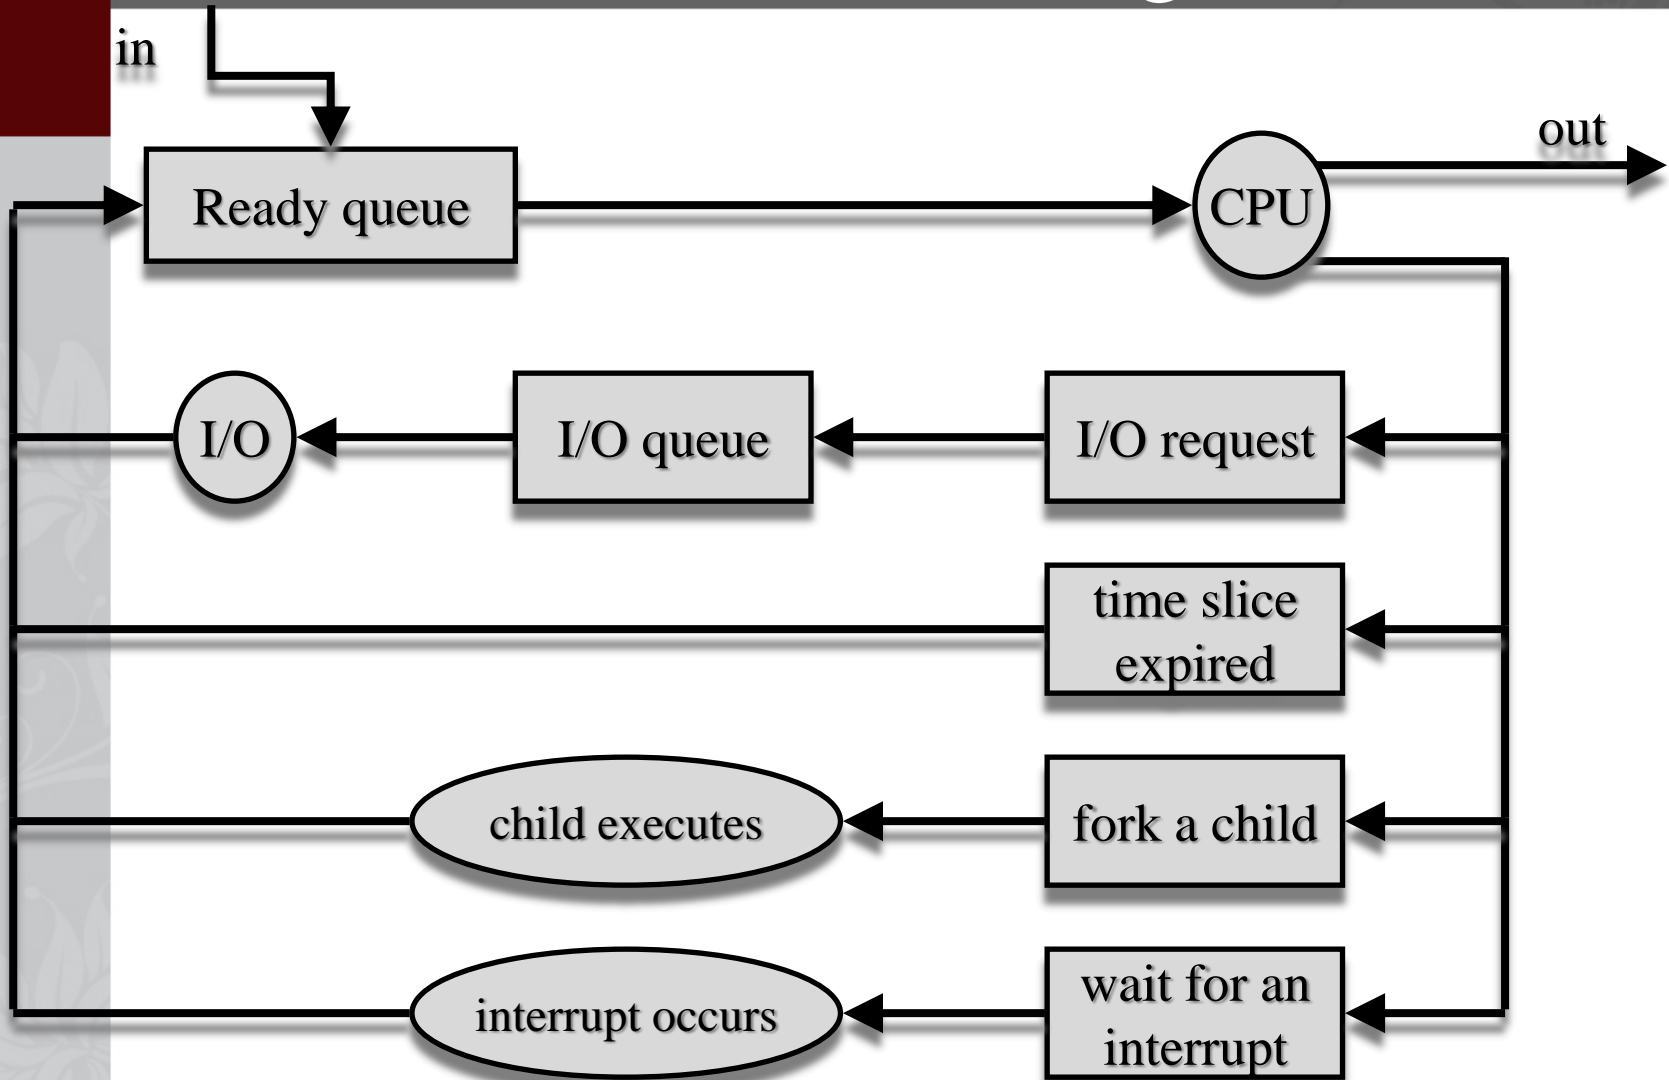
- **Job queue**
 - set of all processes in the system.
- **Ready queue**
 - set of all processes residing in main memory which is ready and waiting to execute.
- **Device queues**
 - set of processes waiting for an I/O device.

* Process migration between the various queues

Ready Queue And Various I/O Device Queues



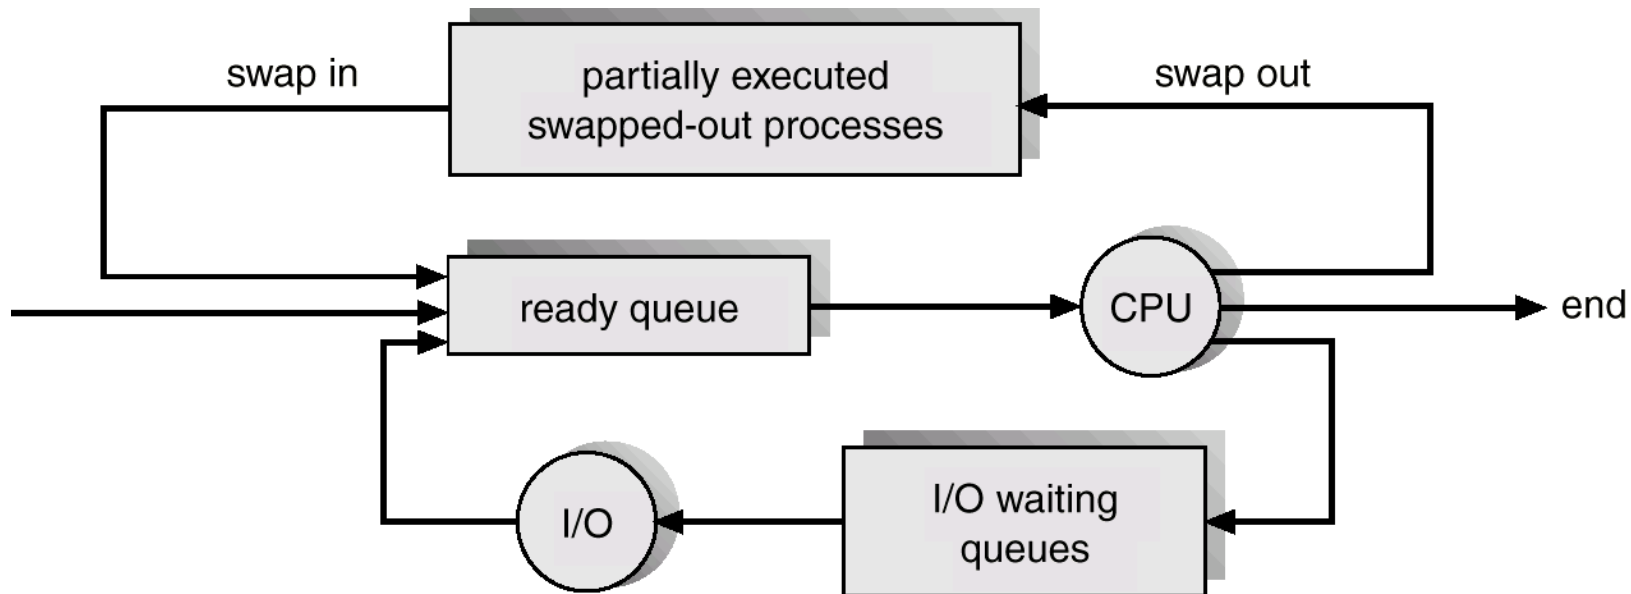
Representation of Process Scheduling



Schedulers

- **Long-term scheduler (or job scheduler)**
 - Selects which processes should be brought into the ready queue.
 - Executes infrequently
 - May be absent in some O/S
- **Short-term scheduler (or CPU scheduler)**
 - Selects which process should be executed next and allocates CPU.
 - Executes frequently

Addition of Medium Term Scheduling



Schedulers Cont'd

- Processes can be described as either:
 - *I/O-bound process*
 - spends more time doing I/O than computations, many short CPU bursts.
 - *CPU-bound process*
 - spends more time doing computations; few very long CPU bursts.
- Proper System performance
 - Mix of CPU & I/O bound processes
- Improper system performance
 - All processes are I/O bound
 - All processes are CPU bound

Context Switch

- When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process.
- Context-switch time is overhead; the system does no useful work while switching.
- Time dependent on hardware support.

OPERATIONS ON PROCESSES

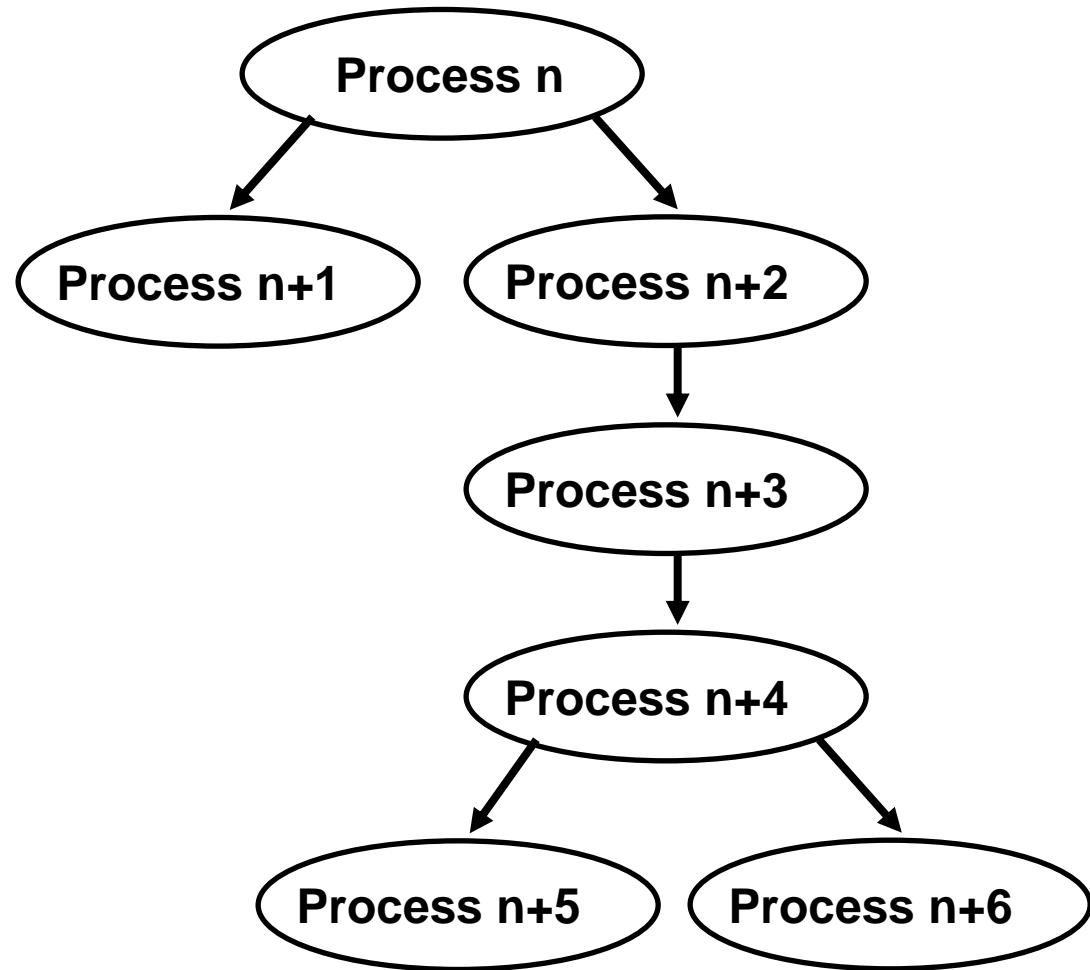

Operations on Processes

- Process Creation
- Process Termination

* Process resources

- CPU time
- Memory
- Files
- I/O devices

Process Creation



Process Creation Cont'd

- Parent process create children processes, which, in turn create other processes, forming a tree of processes.
- Resource sharing
 - Parent and children share all resources.
 - Children share subset of parent's resources.
- Execution
 - Parent and children execute concurrently.
 - Parent waits until children terminate.

Process Creation Cont'd

- **Address space**
 - Child duplicate of parent.
 - Child has a program loaded into it.
- **UNIX examples**
 - **fork** system call creates new process
 - **exec** system call used after a **fork** to replace the process' memory space with a new program.

Process Termination

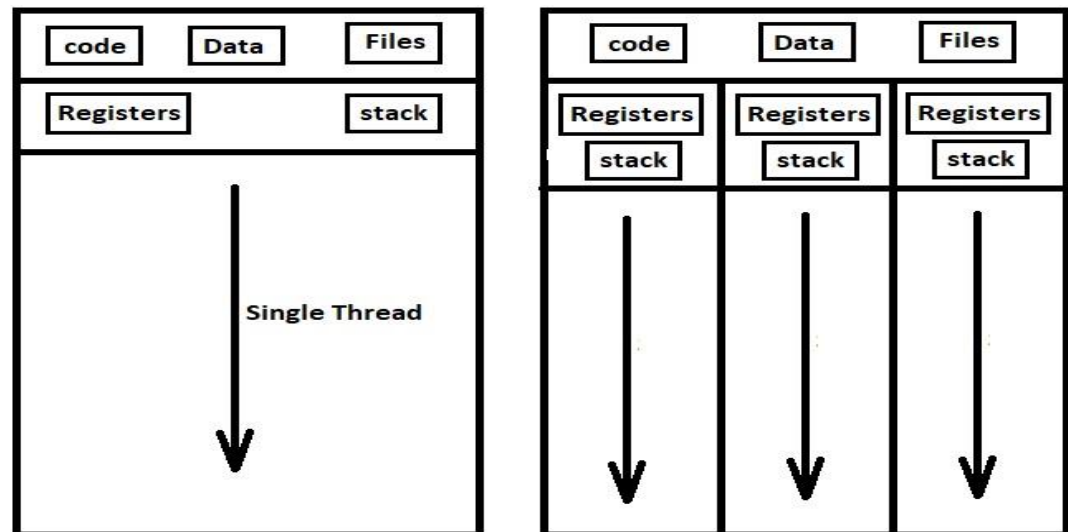
- Process executes last statement and asks the operating system to decide it (**exit**).
 - Output data from child to parent.
 - Process' resources are de-allocated by operating system.
- Parent may terminate execution of children processes (**abort**).
 - Child has exceeded allocated resources.
 - Task assigned to child is no longer required.
 - Parent is exiting.
 - Operating system does not allow child to continue if its parent terminates.
 - Cascading termination.

Threads

- Many modern operating systems now provide features for a process to contain multiple threads of control instead of a single Thread
- A thread, sometimes called a lightweight process (LWP)

Threads

- It shares with other threads belonging to the same process its **code section**, **data section**, and other operating-system resources, such as **open files**



Threads

- **Benefits:**
 1. Responsiveness
 2. Resource sharing
 3. Economy
 4. Utilization of multiprocessor architectures

Multi Threading Models

- **Many-to-One Model:**
 - Maps many user-level threads to one kernel thread. Thread
 - Management is done in user space, so it is efficient
 - The entire process will block if a thread makes a blocking system call
 - Multiple threads are unable to run in parallel on multiprocessors

Multi Threading Models

- **One-to-one Model:**
 - Maps each user thread to a kernel thread.
 - It provides more concurrency than the many-to-one model.
 -
 - Another thread can run when a thread makes a blocking system call.
 -
 - It also allows multiple threads to run in parallel on multiprocessors.
 - The only drawback is the overhead of creating kernel threads can burden the performance of an application.

Multi Threading Models

- **Many-to-Many Model:**
 - Multiplexes many user-level threads to a smaller or equal number of kernel threads.
 - Developers can create as many user threads as necessary, and the corresponding kernel threads can run in parallel on a multiprocessor.
 - Also, when a thread performs a blocking system call, the kernel can schedule another thread for execution.

Threading Issues

- **The fork and exec System Calls**
 - Does the new process duplicate all threads or is the new process single-threaded?
- **Cancellation**
 - Allocating resources to a cancelled thread or cancel a thread while updating data that is shared by another thread
- **Signal Handling**
 - Handling signals in single-threaded programs is straightforward; signals are always delivered to a process. However, delivering signals is more complicated in multithreaded programs, as a process may have several threads. Where then should a signal be delivered?
- **Thread Pools**
 - Unlimited threads could exhaust system resources, such as CPU time or memory. One solution to this issue is to use thread pools
- **Thread-Specific Data**
 - Each thread might need its own copy of certain data in some circumstances.

Cooperating Processes

- *Independent* process cannot affect or be affected by the execution of another process.
- *Cooperating* process can affect or be affected by the execution of another process
- Advantages of process cooperation
 - Information sharing
 - Computation speed-up
 - Modularity
 - Convenience

INTER-PROCESS COMMUNICATION

Inter-process Communication

- Message Passing
- Shared Memory
- * Synchronization

Message Passing

- Messages
 - Fixed size
 - Variable size
- Communication
 - Direct:
 - `send(P2, message), receive(P1, message)`
 - Indirect:
 - `send(ID, message), receive(ID, messages)`

Synchronization

- Message passing may be either blocking or non-blocking.
- **Blocking** is considered **synchronous**
- **Non-blocking** is considered **asynchronous**
- **send** and **receive** primitives may be either blocking or non-blocking.

